Can Silicon Be Conscious? A TNA Analysis of the Limits of Operational Systems

Author: Claudio Bresciano

Abstract

The rapid development of artificial intelligence has revived one of the oldest philosophical questions: can a purely physical machine become conscious? Contemporary materialist theories generally assume that consciousness emerges from sufficiently complex information processing, regardless of the substrate in which that processing occurs. Under this view, silicon and biological neurons differ only in implementation, not in principle.

This paper argues that such a conclusion remains unproven. Drawing upon the Theory of Axiomatic Necessity (TNA), the Kripke–Wittgenstein rule-following paradox, and the theorem of Failure of Local Closure, we show that operational dynamics alone may be insufficient to derive semantic content, normativity, and subjective experience. If consciousness requires more than computation, then increasing computational complexity cannot by itself guarantee the emergence of awareness.

The question is therefore not whether silicon can simulate consciousness, but whether consciousness is reducible to operational states in the first place.

1. Introduction

The dominant assumption behind contemporary artificial intelligence is simple:

If a machine performs the same functions as a conscious human mind, then the machine itself is conscious.

This assumption underlies much of modern functionalism and computational theories of mind.

Yet the inference remains controversial.

A calculator may perfectly reproduce arithmetic operations without understanding arithmetic.

Likewise, a language model may reproduce meaningful sentences without possessing meaning.

The central question is therefore:

> *Can operational equivalence produce consciousness?*

Or, more fundamentally:

> *Is consciousness derivable from operational dynamics alone?*

TNA approaches this question through the distinction between operational systems (N_0) and admissibility structures (N_1).

2. The Materialist Hypothesis

The standard materialist argument proceeds as follows:

1. Consciousness occurs in biological brains.
2. Brains are physical systems.
3. Therefore consciousness is generated by physical processes.
4. Any sufficiently similar physical process should generate consciousness.

Under this view, the substrate is irrelevant.

Carbon-based neurons and silicon-based transistors differ only in implementation.

Consciousness is assumed to emerge whenever the correct computational organization is achieved.

The difficulty is that the conclusion depends upon a premise that remains unverified:

> *That consciousness is fully reducible to operational structure.*

This is precisely the point under dispute.

3. The Problem of Meaning

Consider Kripke's famous distinction between addition and quaddition.

A finite operational history is compatible with multiple interpretations.

The system's behavior alone does not determine the rule.

The same problem appears in artificial intelligence.

A language model may generate the sentence:

"Montevideo is the capital of Uruguay."

The operational description consists of:

- numerical weights,
- matrix multiplications,
- activation functions,
- probability distributions.

Yet none of these states contains "Montevideo" as experienced meaning.

The machine produces symbols.

Whether those symbols possess semantic content remains an independent question.

The problem is therefore not computational competence but semantic grounding.

4. The Problem of Experience

The challenge becomes even stronger when subjective experience is considered.

Suppose a machine perfectly identifies colors.

It labels red correctly.

It labels blue correctly.

Its outputs are indistinguishable from those of a human observer.

Does this imply that the machine experiences red?

Materialism typically assumes that sufficient complexity will eventually generate experience.

Yet this assumption faces the classical challenge of qualia.

Red and blue are not merely different information states.

They are different experiences.

The operational description may distinguish two states:

State A

State B

But the existence of a phenomenal difference between them is not obviously derivable from the operational description itself.

The problem is not classification.

The problem is experience.

5. Failure of Local Closure

The TNA theorem of Failure of Local Closure states:

> *A local operational system cannot generate the conditions that legitimize its own interpretation.*

Applied to consciousness:

N_0 consists of:

- neural activity,
- silicon circuits,
- computational states,
- information processing.

The question is whether N_0 can generate:

- meaning,
- normativity,
- subjective experience.

Materialism assumes that it can.

TNA argues that this remains unproven.

Increasing the complexity of N_0 does not automatically demonstrate the emergence of N_1.

Adding more operations does not necessarily generate admissibility.

6. The Silicon Argument

Advocates of artificial consciousness often argue:

"If neurons can produce consciousness, why not silicon?"

The TNA response is:

"If consciousness is purely operational, then silicon should indeed be sufficient."

However, this merely shifts the problem.

The real issue is not the material substrate.

The real issue is whether consciousness is operational at all.

Replacing neurons with transistors does not solve the explanatory gap.

It simply relocates it.

The question remains:

Why should any physical process be accompanied by subjective experience?

7. The Dark Matter Analogy

Contemporary discussions of consciousness sometimes resemble cosmological appeals to unseen explanatory variables.

When a phenomenon cannot be derived from existing theory, additional mechanisms are proposed.

The materialist expectation is that future neuroscience will eventually explain consciousness.

This may be correct.

But it remains a prediction, not a demonstration.

TNA suggests that the explanatory difficulty may be structural rather than merely incomplete.

If consciousness belongs partly to N_1 rather than entirely to N_0, then no amount of additional operational detail will eliminate the gap.

8. Conclusion

The question "Can silicon be conscious?" cannot be answered by measuring computational power alone.

Before asking whether silicon can host consciousness, one must first answer a deeper question:

> *Is consciousness reducible to operational dynamics?*

The Kripke–Wittgenstein paradox demonstrates that operational behavior does not uniquely determine meaning.

The problem of qualia suggests that operational states do not obviously determine experience.

The theorem of Failure of Local Closure suggests that admissibility cannot be derived entirely from the domain it governs.

Consequently, the possibility of conscious silicon remains undecidable within purely operational frameworks.

The debate is not ultimately about silicon.

It is about whether consciousness belongs entirely to N_0, or whether every conscious system necessarily presupposes an irreducible N_1.

Until that question is resolved, artificial consciousness remains a philosophical hypothesis rather than an established scientific fact.

References

1. Kripke, S. A. (1982). *Wittgenstein on Rules and Private Language*. Harvard University Press.
2. Chalmers, D. J. (1996). *The Conscious Mind*. Oxford University Press.
3. Bresciano, C. (2026). *Failure of Local Closure in Self-Optimizing Systems: A Minimal Structural Bound*. Zenodo.
4. Bresciano, C. (2025). *Core Structure of the Theory of Axiomatic Necessity*. Zenodo.
5. Searle, J. R. (1980). "Minds, Brains and Programs." *Behavioral and Brain Sciences*.

6. Nagel, T. (1974). "What Is It Like to Be a Bat?" *Philosophical Review*.